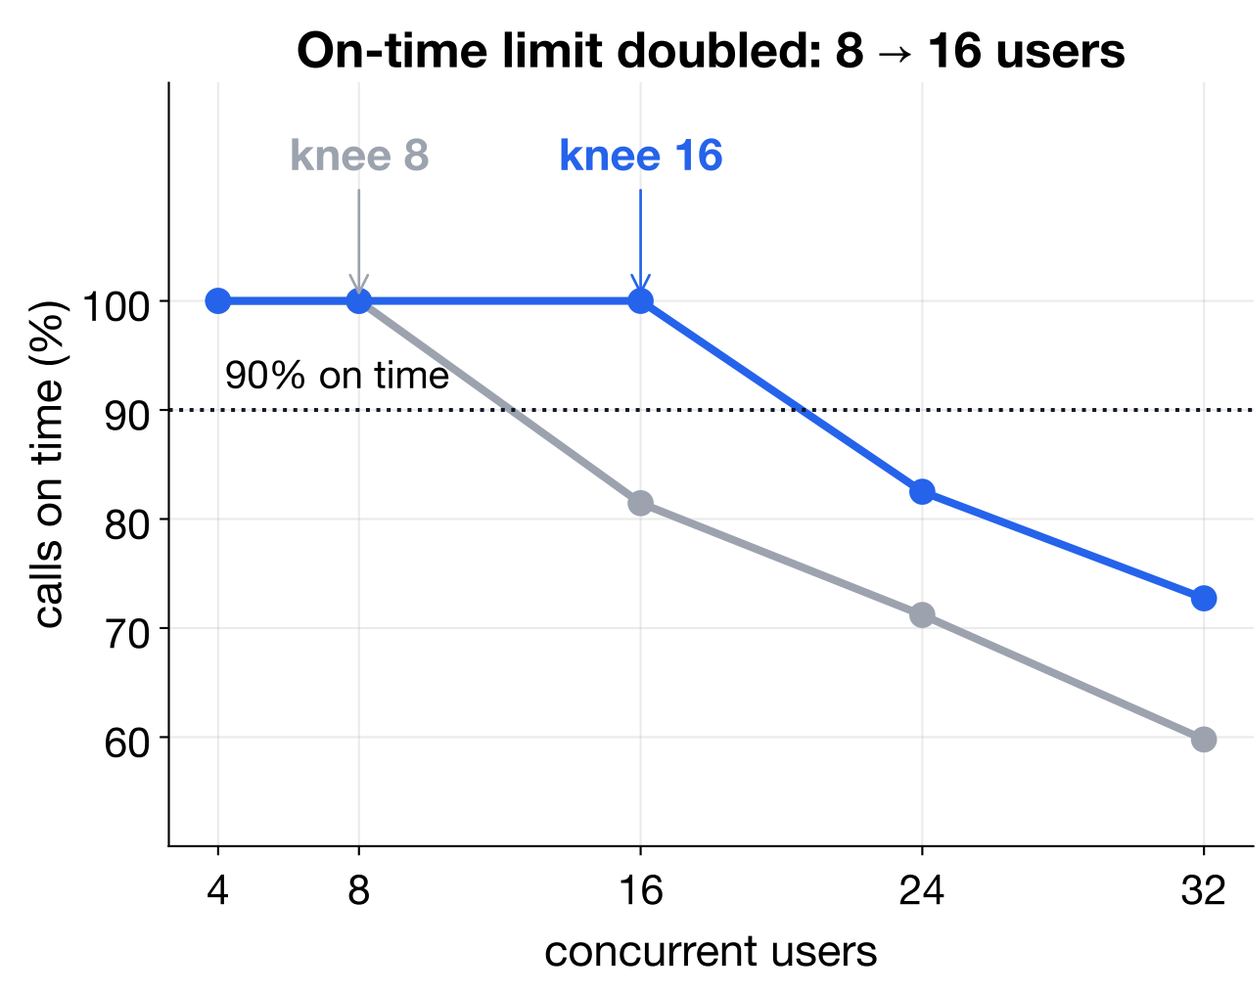
On-time limit doubled: 8 → 16 users
knee 8
knee 16
90% on time
calls on time (%)
100
90
80
70
60
4
8
16
24
32
concurrent users

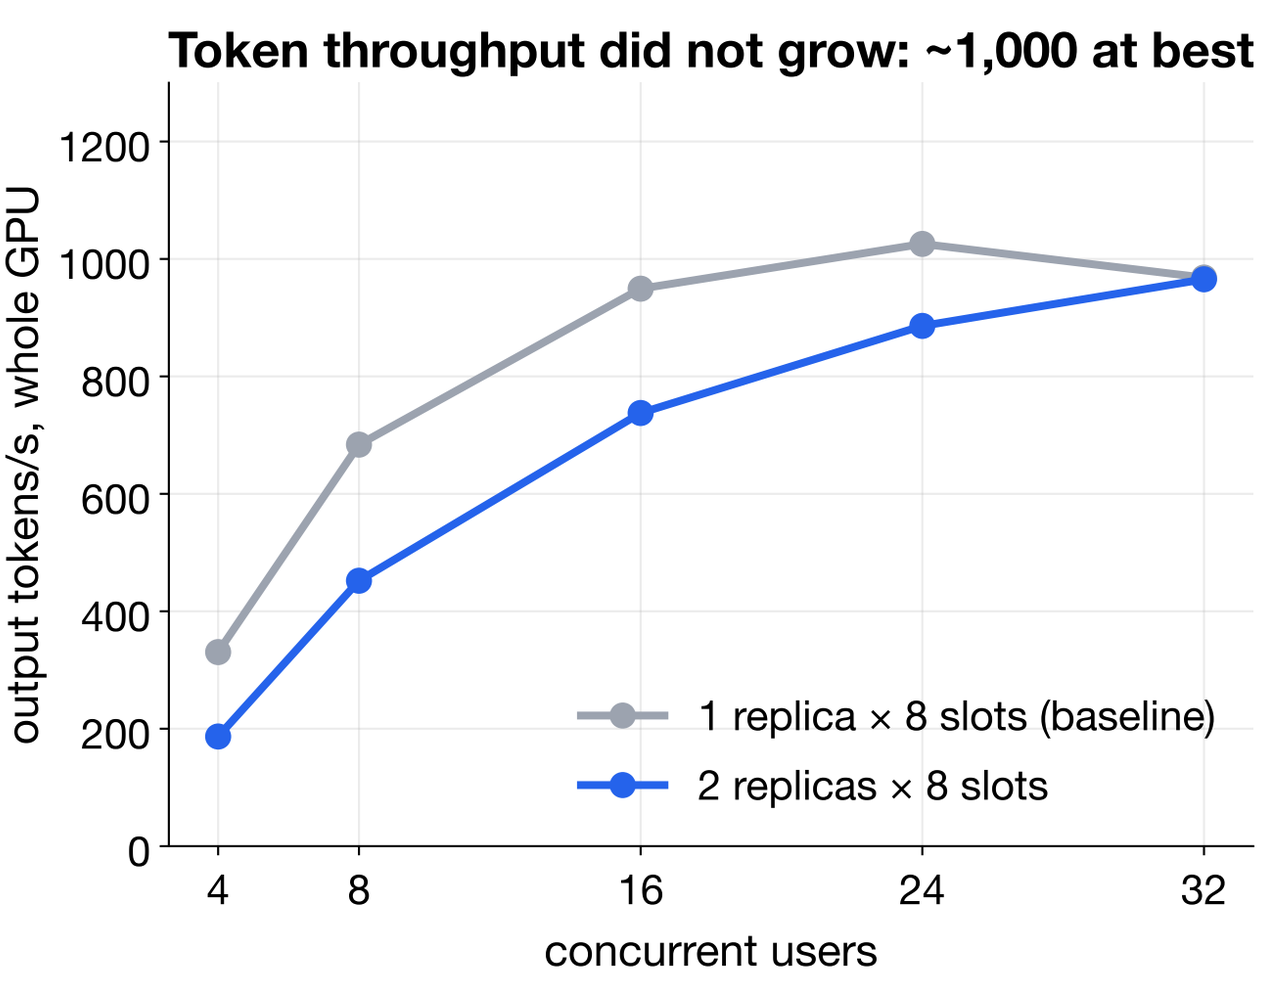
Token throughput did not grow: ~1,000 at best
output tokens/s, whole GPU
1200
1000
800
600
400
200
0
1 replica × 8 slots (baseline)
2 replicas × 8 slots
4
8
16
24
32
concurrent users