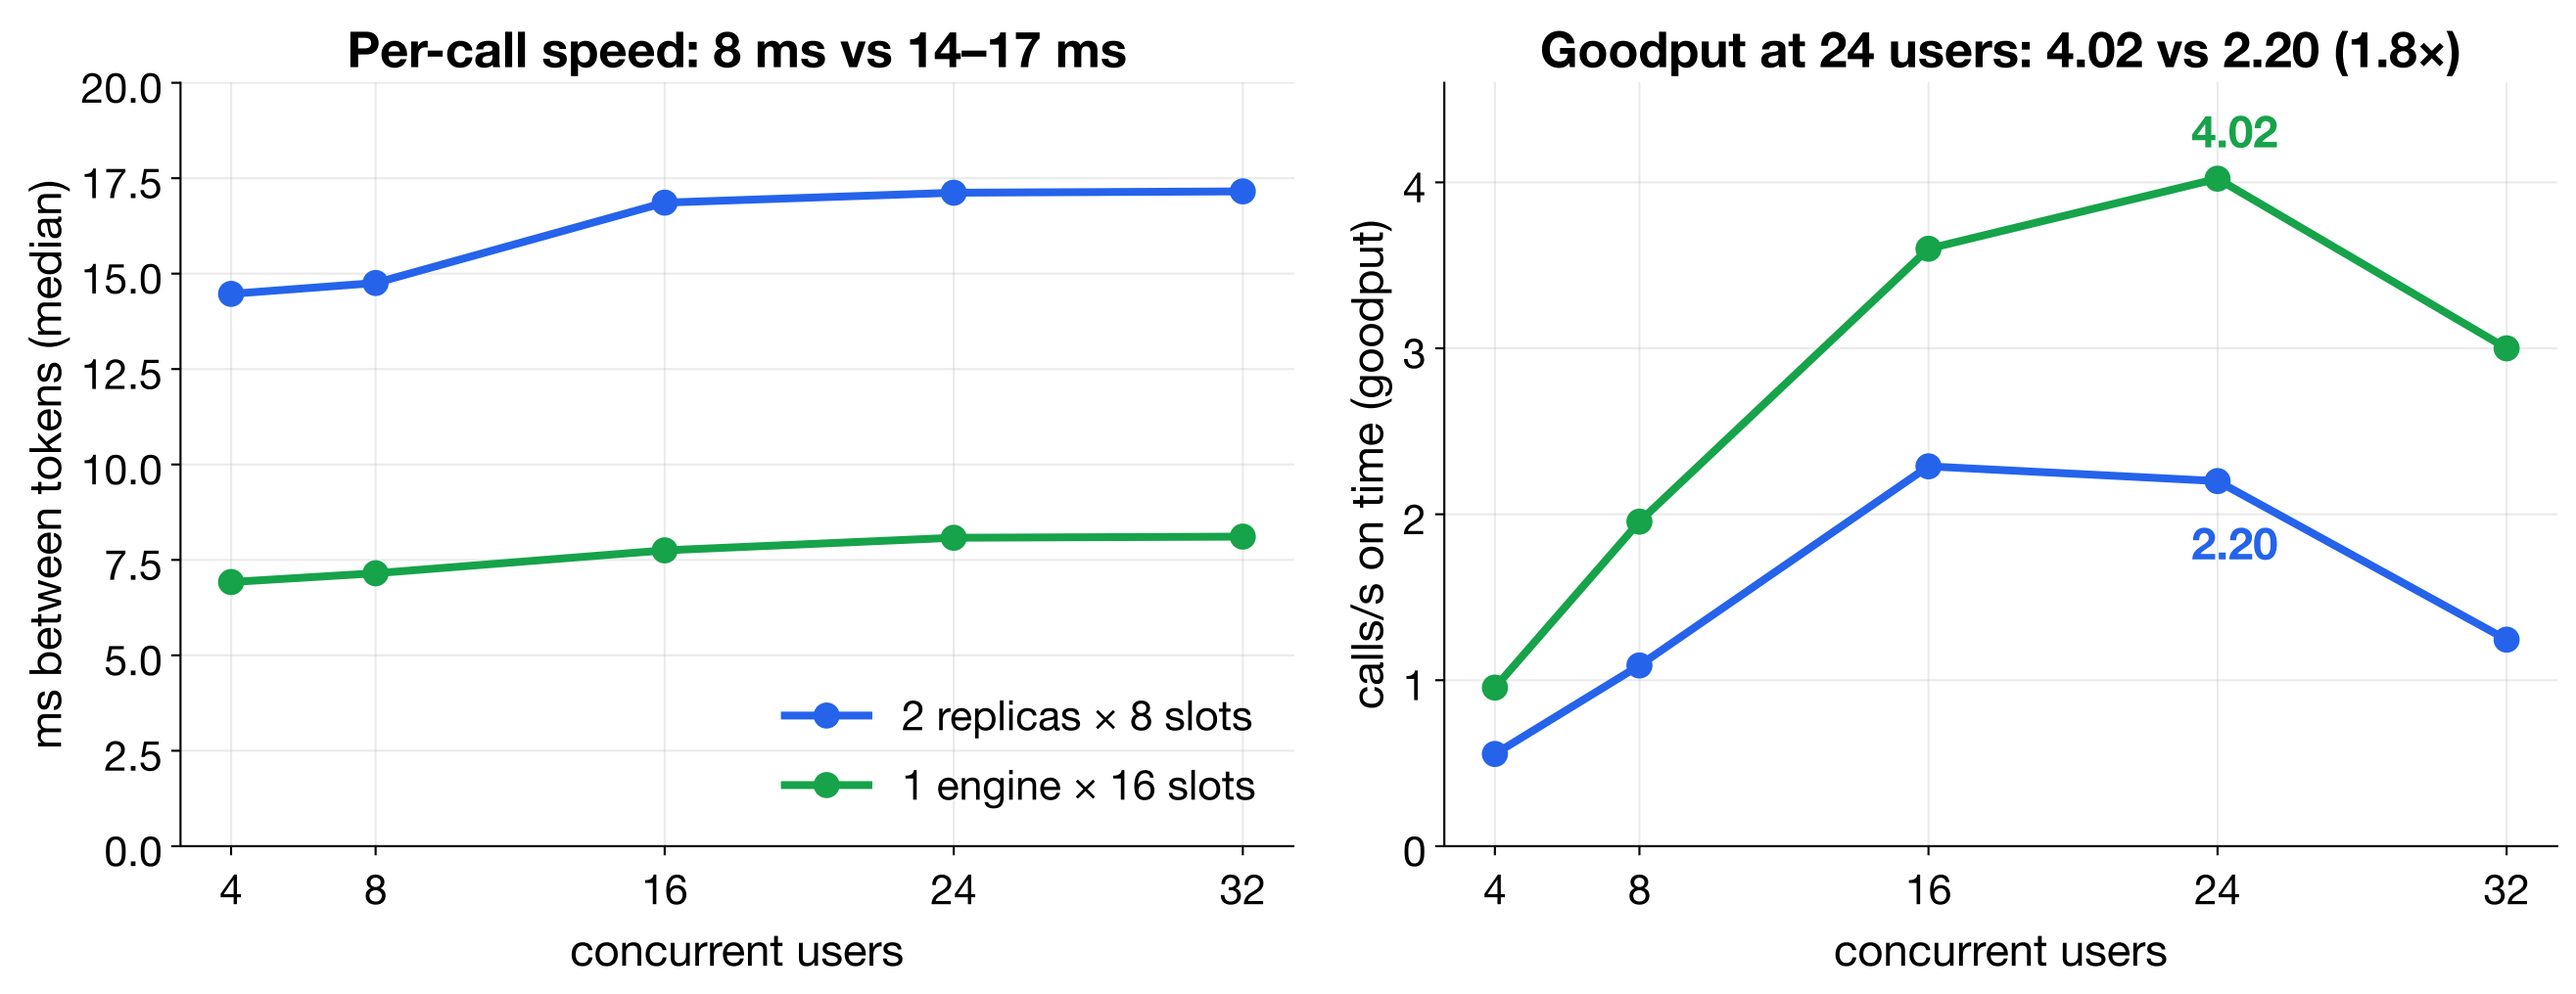
Per-call speed: 8 ms vs 14–17 ms
ms between tokens (median)
20.0
17.5
15.0
12.5
10.0
7.5
5.0
2.5
0.0
4
8
16
24
32
concurrent users
2 replicas × 8 slots
1 engine × 16 slots
Goodput at 24 users: 4.02 vs 2.20 (1.8×)
calls/s on time (goodput)
4
3
2
1
0
4.02
2.20
4
8
16
24
32
concurrent users