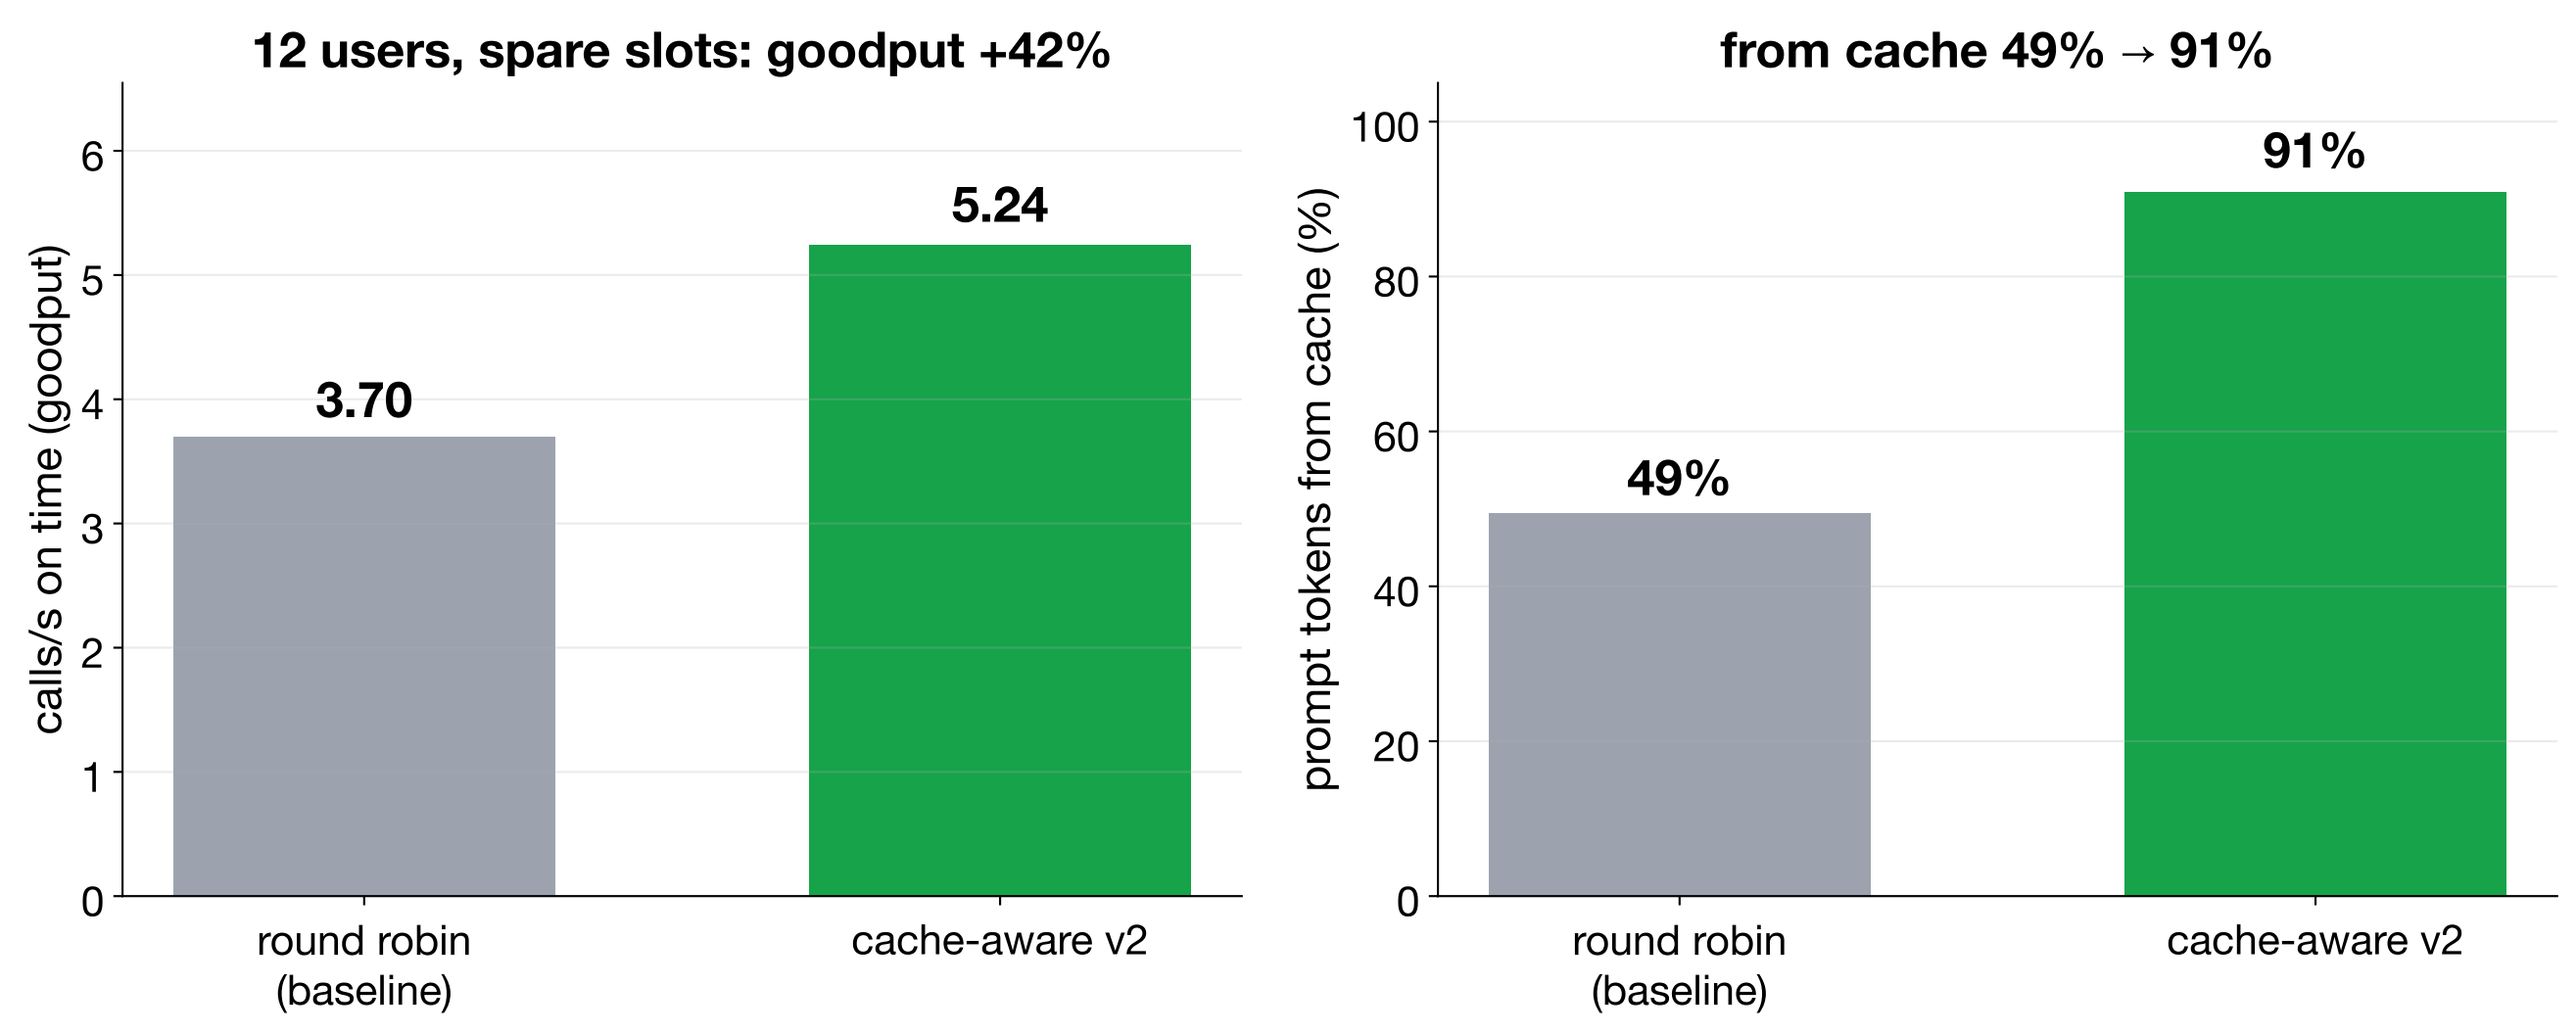
12 users, spare slots: goodput +42%
calls/s on time (goodput)
0
1
2
3
4
5
6
3.70
5.24
round robin
(baseline)
cache-aware v2
from cache 49% → 91%
prompt tokens from cache (%)
0
20
40
60
80
100
49%
91%
round robin
(baseline)
cache-aware v2