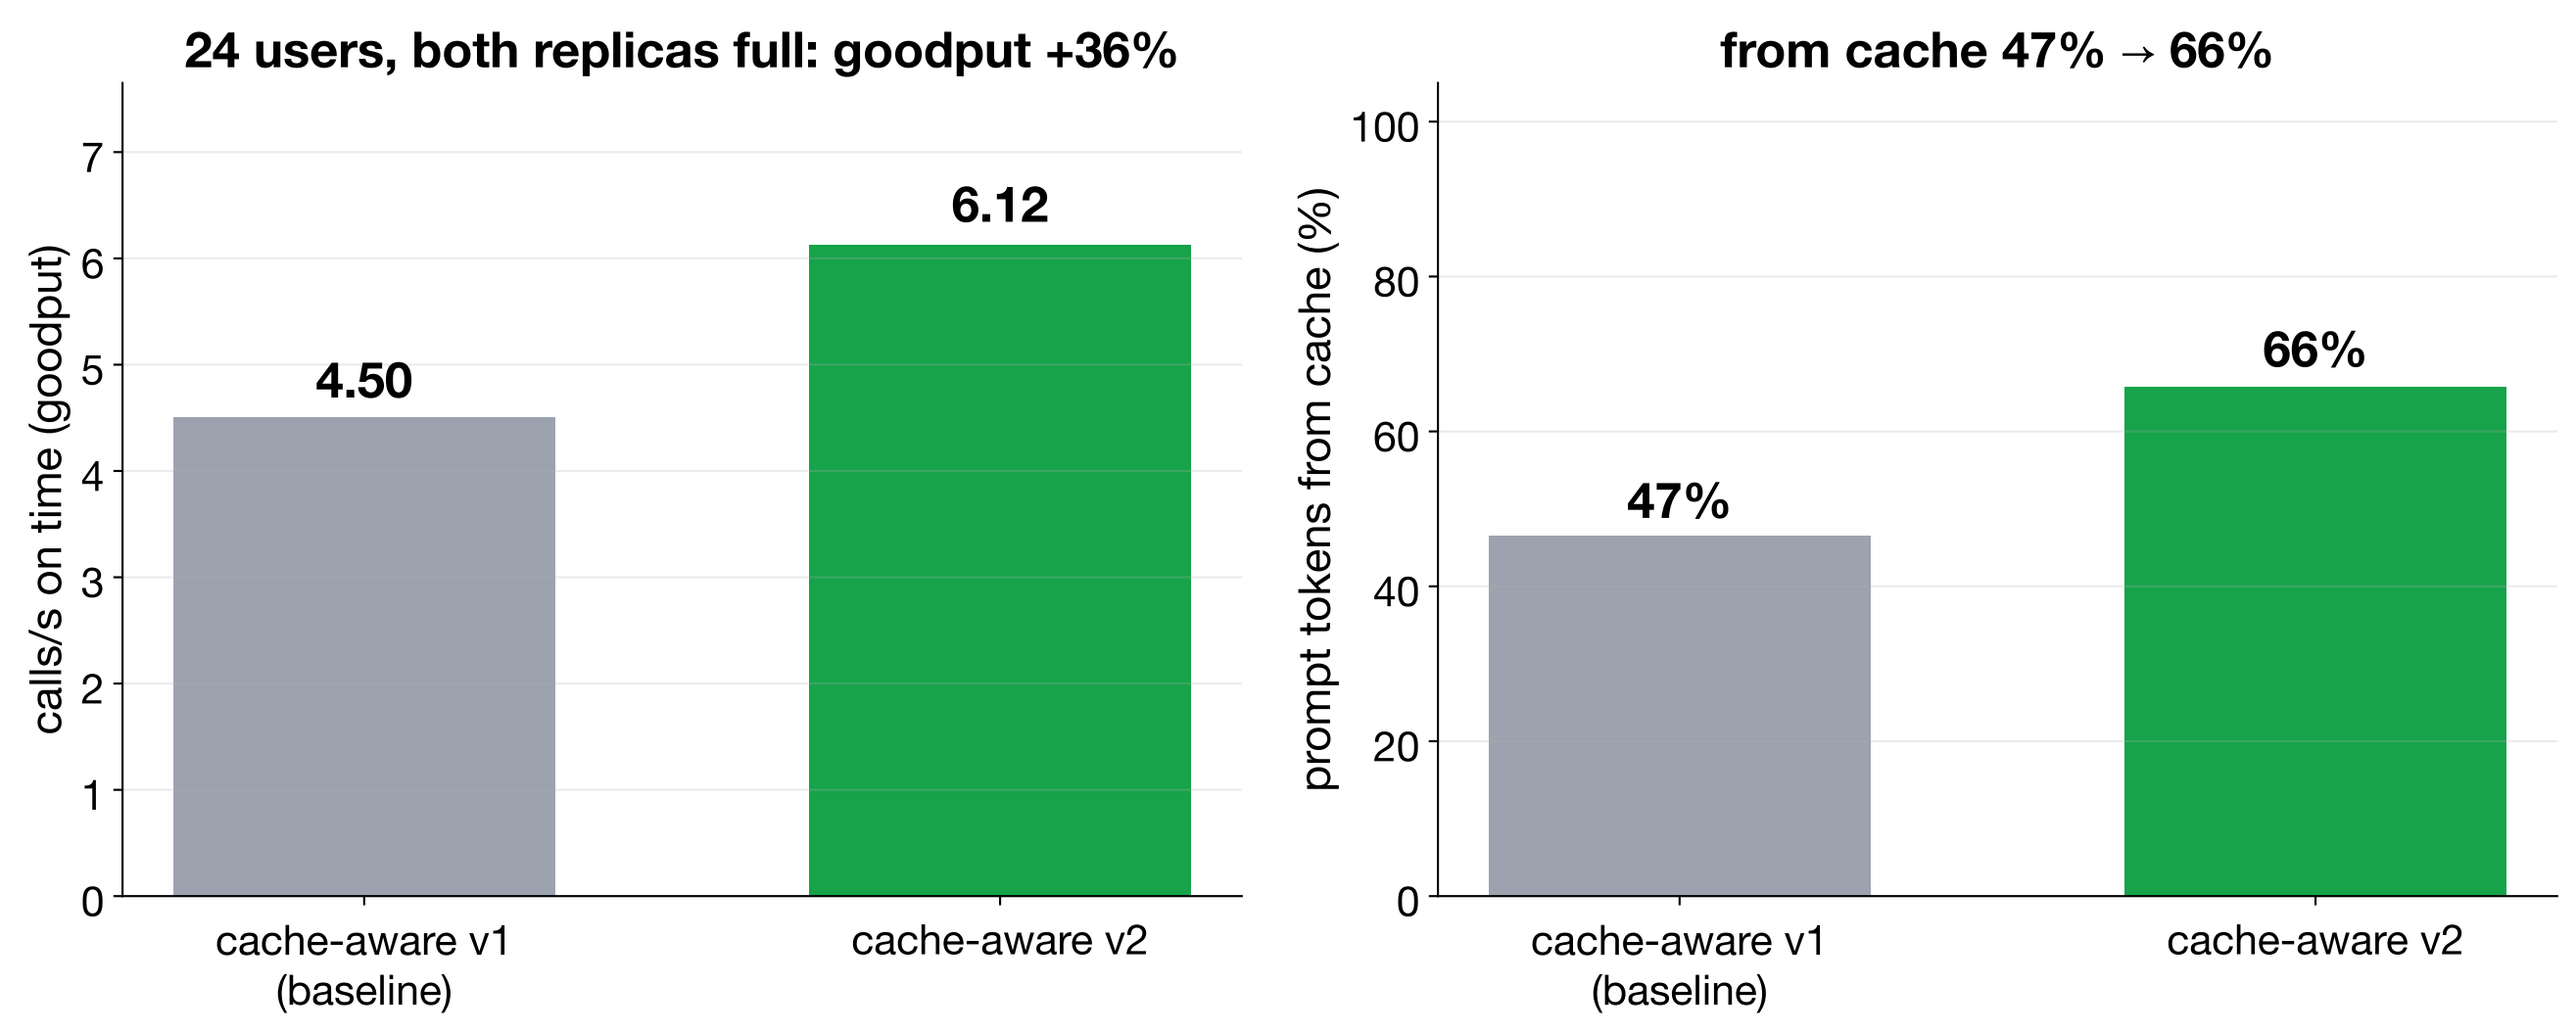
24 users, both replicas full: goodput +36%
calls/s on time (goodput)
0
1
2
3
4
5
6
7
4.50
6.12
cache-aware v1
(baseline)
cache-aware v2
from cache 47% → 66%
prompt tokens from cache (%)
0
20
40
60
80
100
47%
66%
cache-aware v1
(baseline)
cache-aware v2