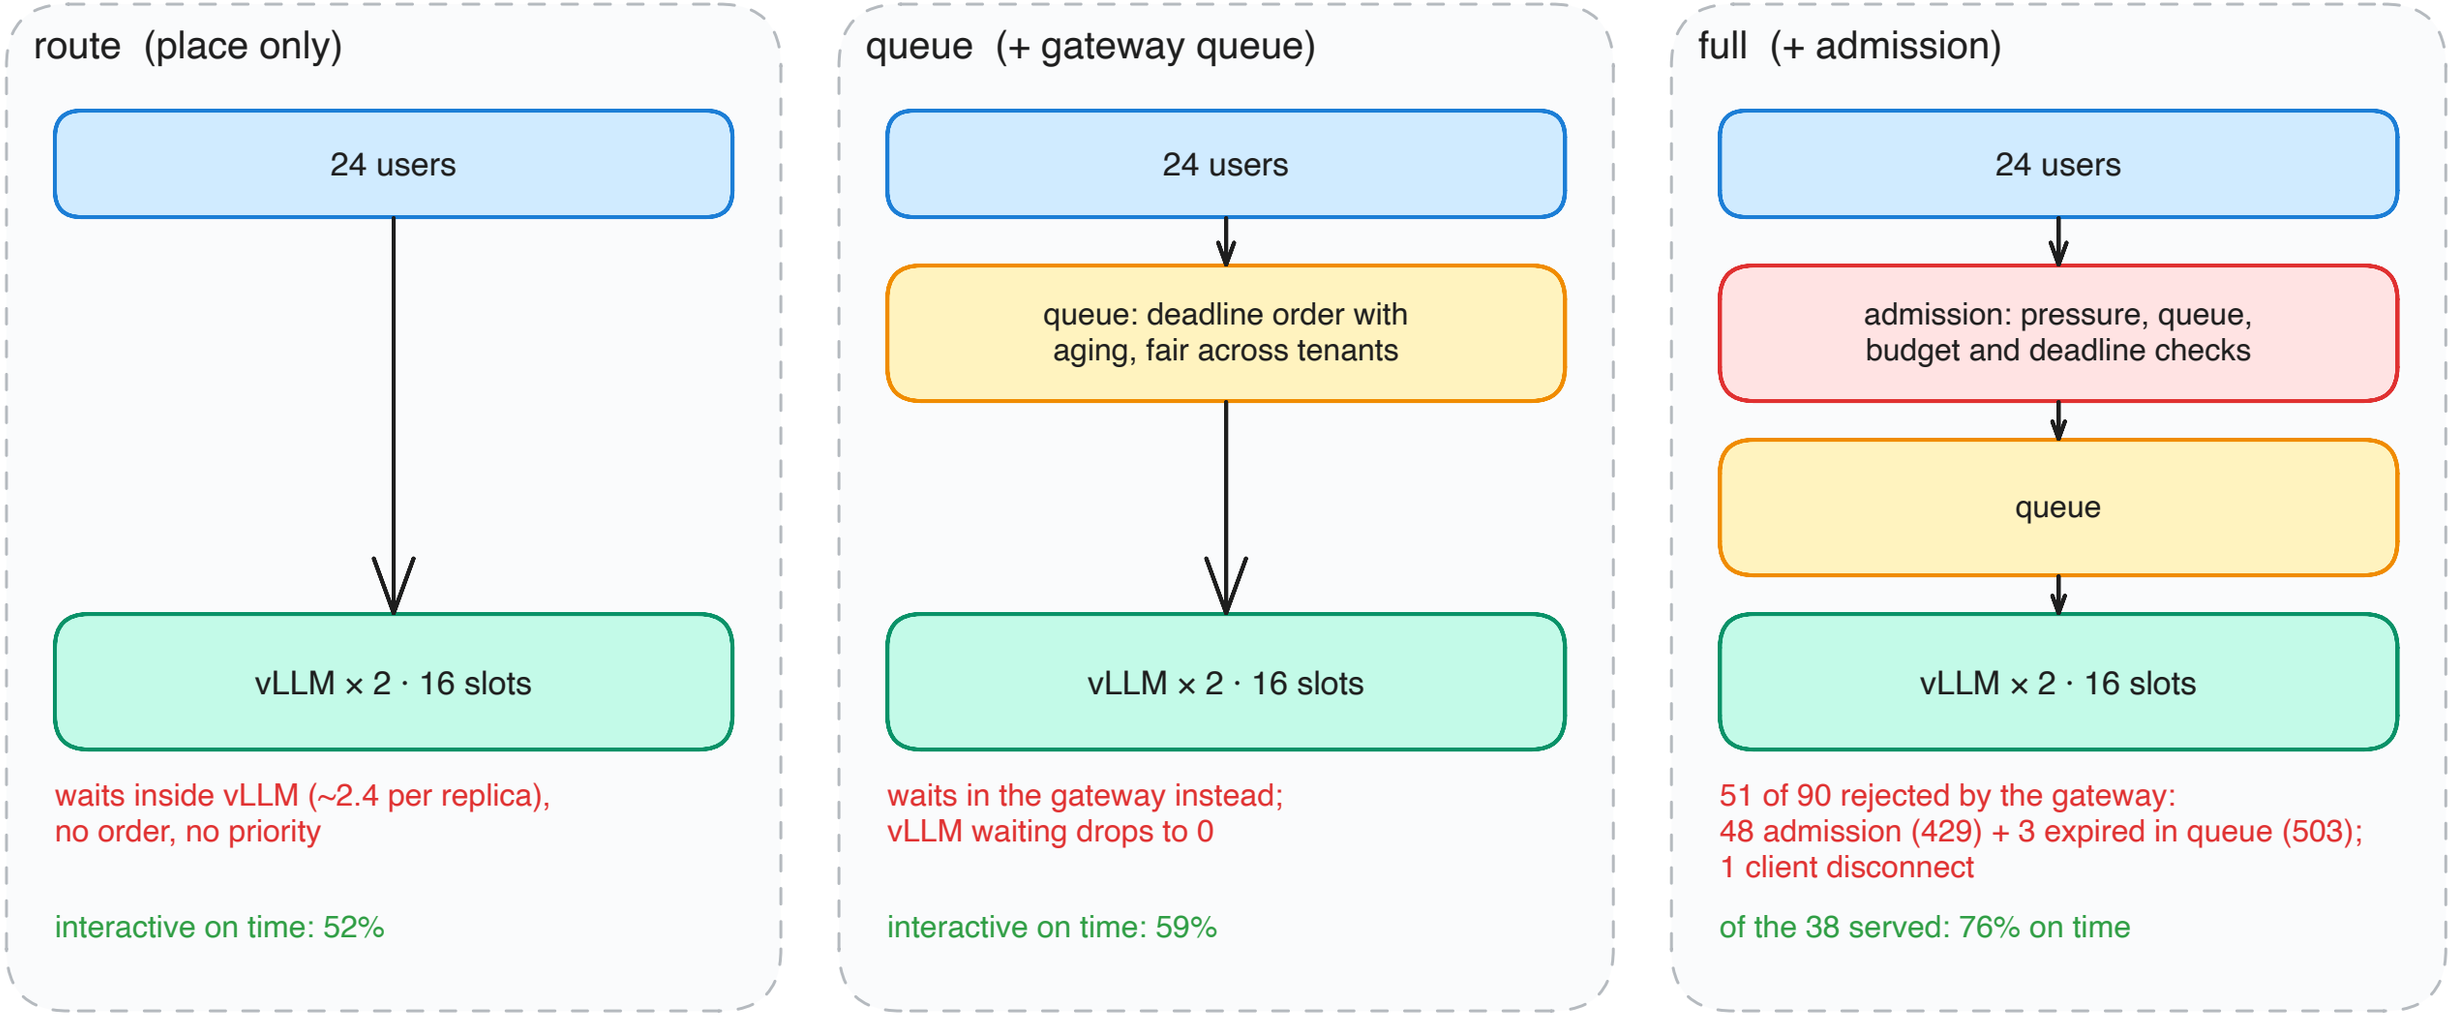
# Where does the overload wait?

24 users on 16 slots: 8 calls too many at any moment. Each gateway mode puts that excess somewhere else (interactive calls).

users · gateway queue · admission · vLLM engines




On time = first token within 2 s. Queue depth is its own experiment (next diagram).