How deep should the gateway queue be?

Mode queue, 24 users on 16 slots. Only the queue bound (per model lane) changes. Time = TTFT p99.

one place in the queue · vLLM engines · many refused · few or none refused

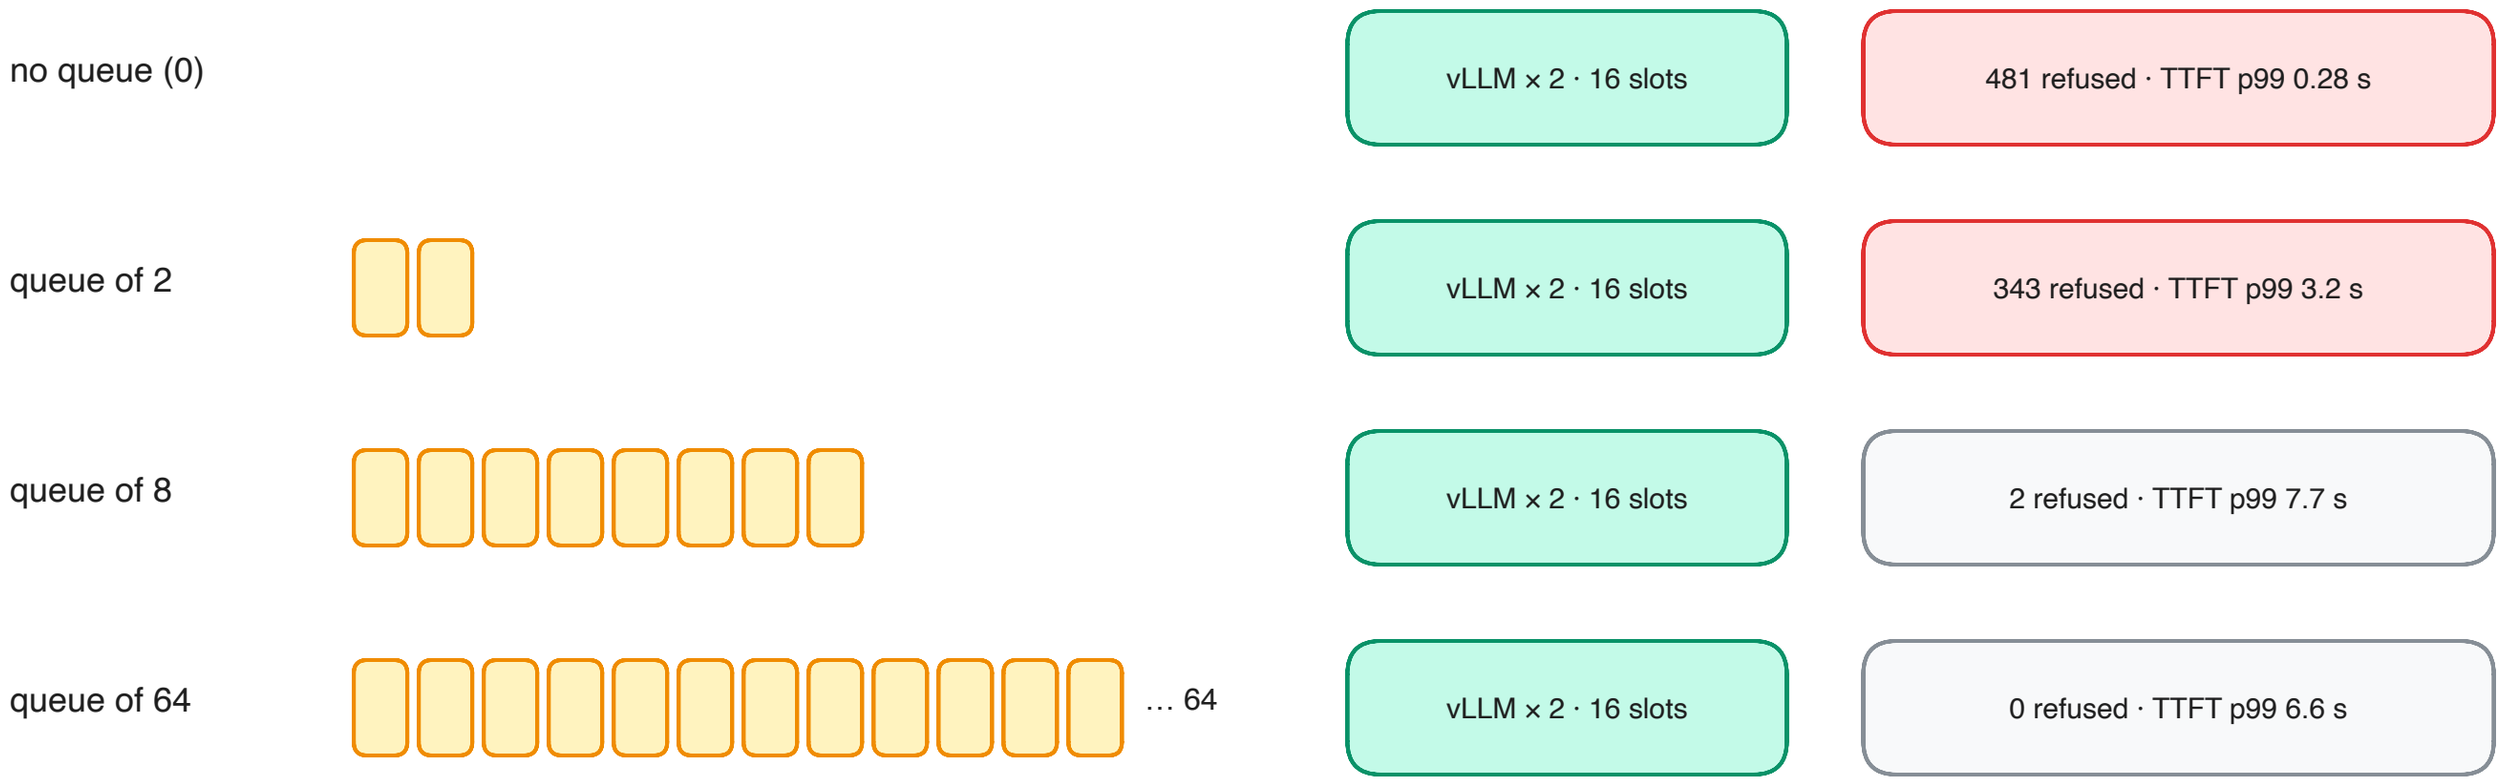

Calls served per second: 2.27 · 2.21 · 2.33 · 2.25 — the same at every depth. Capacity is the GPU's;
a queue decides whether the excess is refused now or waits. (Bound 64's p99 was a little lower than bound 8's: one run each.)