Agent teams run live, by hand

Three tinyharness terminals (SQL, 10-K, code), all agents on the 9B, after the harness fixes

agent teams · gateway · vLLM engines · issue · worked

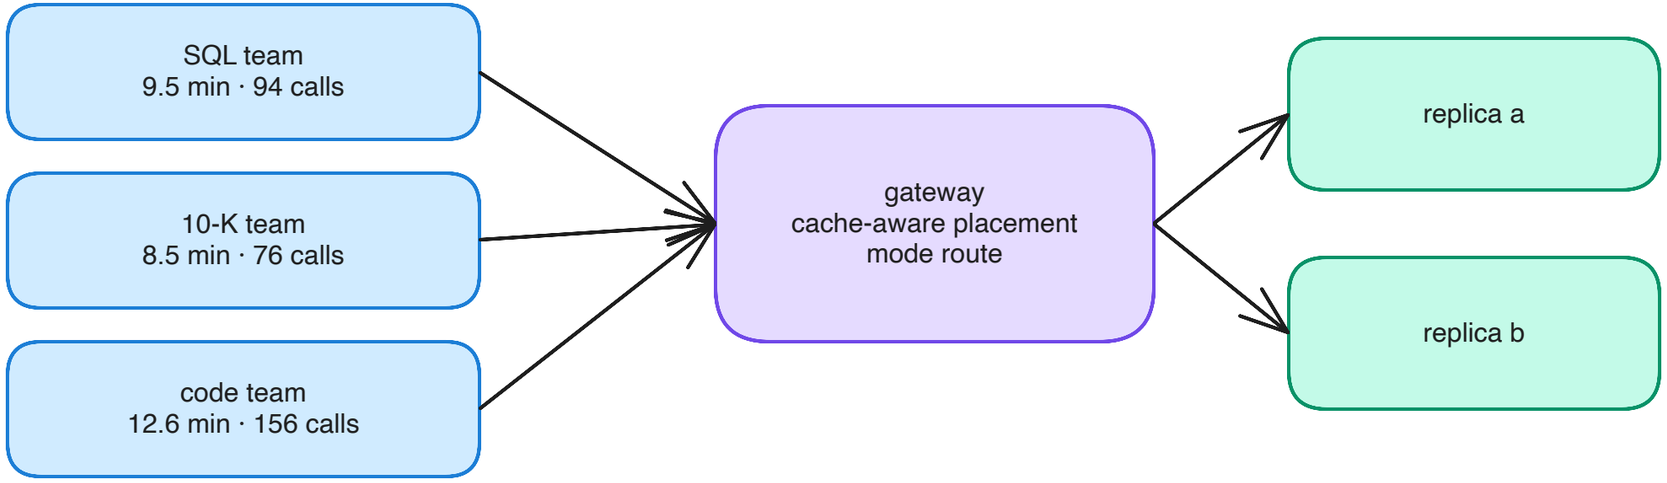

During a brief accidental switch to full (~2 min):
47 admission rejects (429). Every 9B subagent was held to a 3 s first-answer target meant for readers

No sandbox prompts, no team re-runs: the harness fixes held. All three teams completed after returning to route mode.

Report quality fell with the 9B as lead: SQL facts 12% (27B: 64%), 10-K 0 citations (27B: all), code 14 of 16.
That is a model-size finding, not a serving one.