Cache-aware placement with room to choose (12 users)

Same 40 documents, but replicas usually have a free slot, so remembered prefix affinity can hold

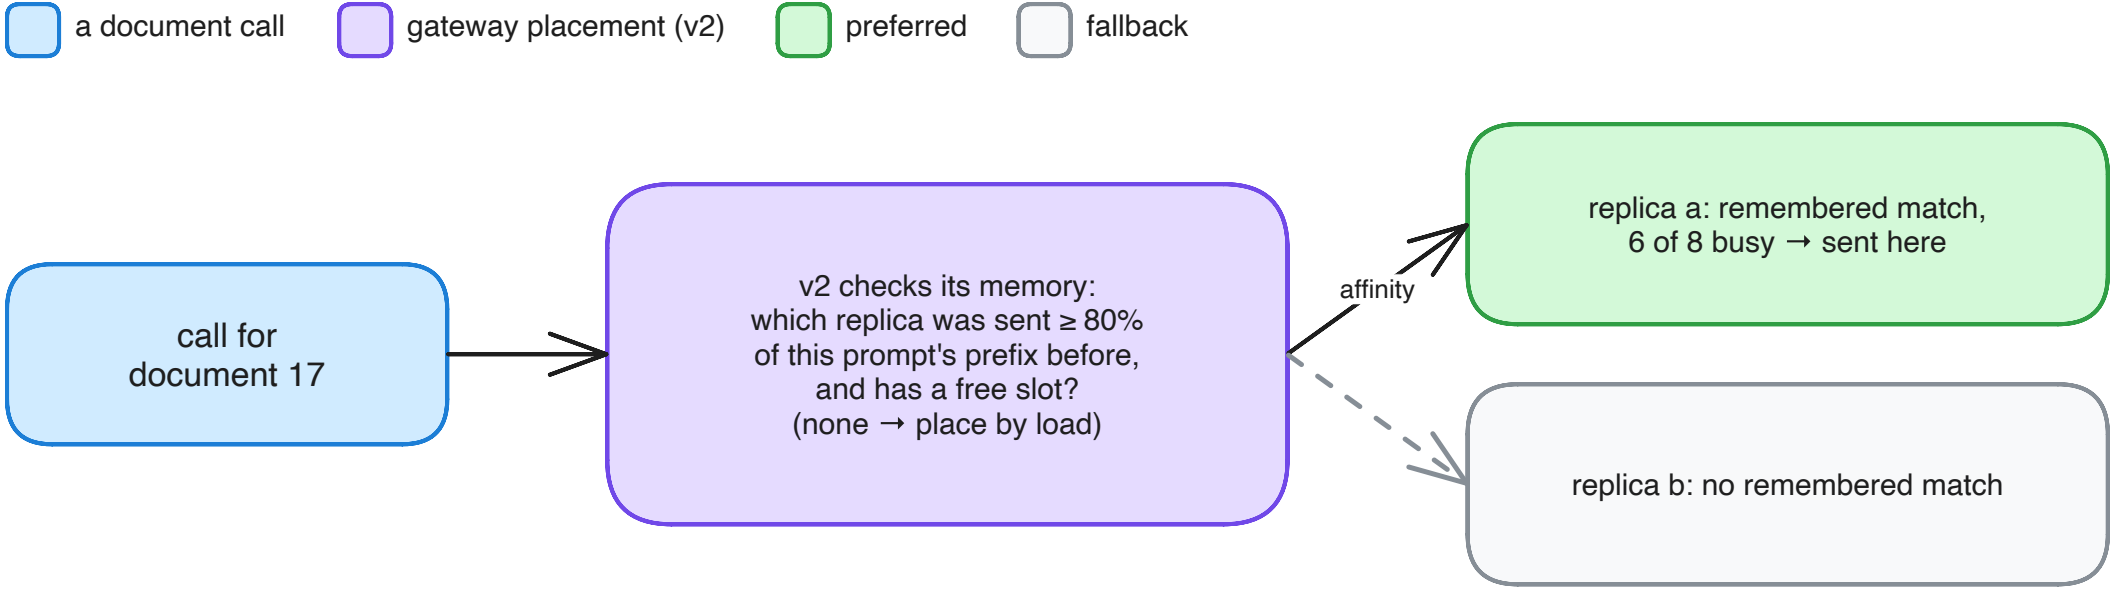

Measured: 86% of each document's calls stayed on one replica, 91% of prompt tokens came from cache (round robin: 49%), first token in 0.10 s instead of 0.62 s, goodput +42% (5.24 vs 3.70 calls/s).