Two bugs in the first cache-aware policy, and the fix

At 24 users both replicas are busy; how the policy breaks ties decides whether the cache helps

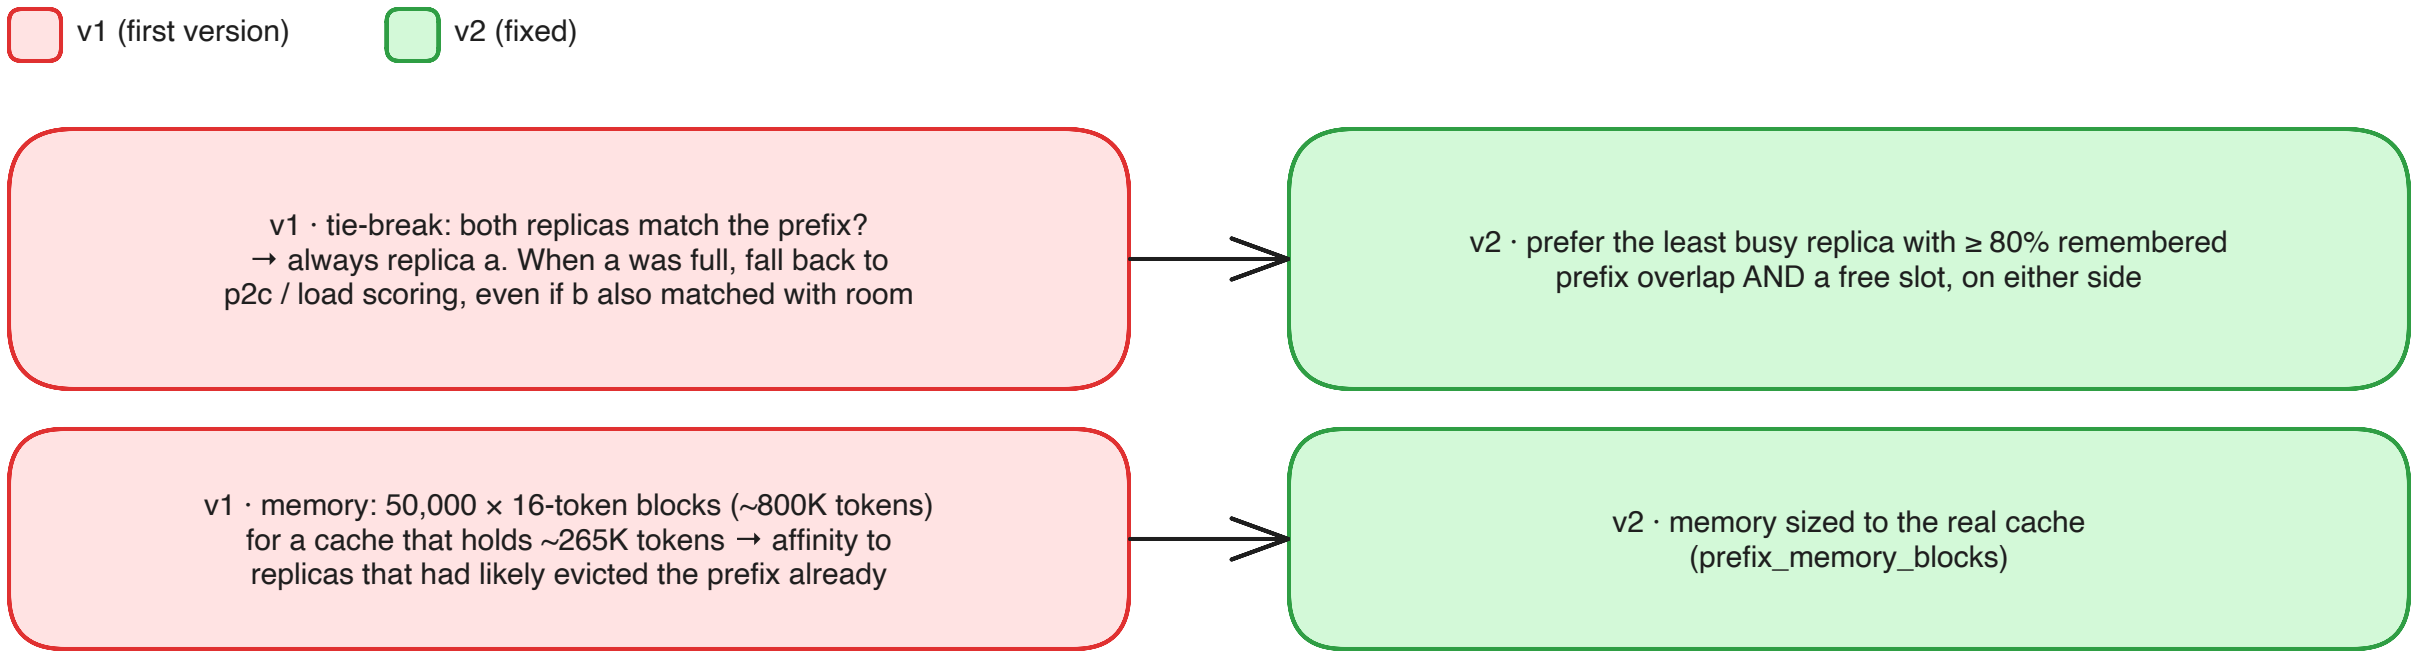

Same load, same documents: goodput 4.50 → 6.12 calls/s (+36%), prompt from cache 47% → 66%.
So Part 2's small gain was mostly v1's bugs, not just the busy replicas.