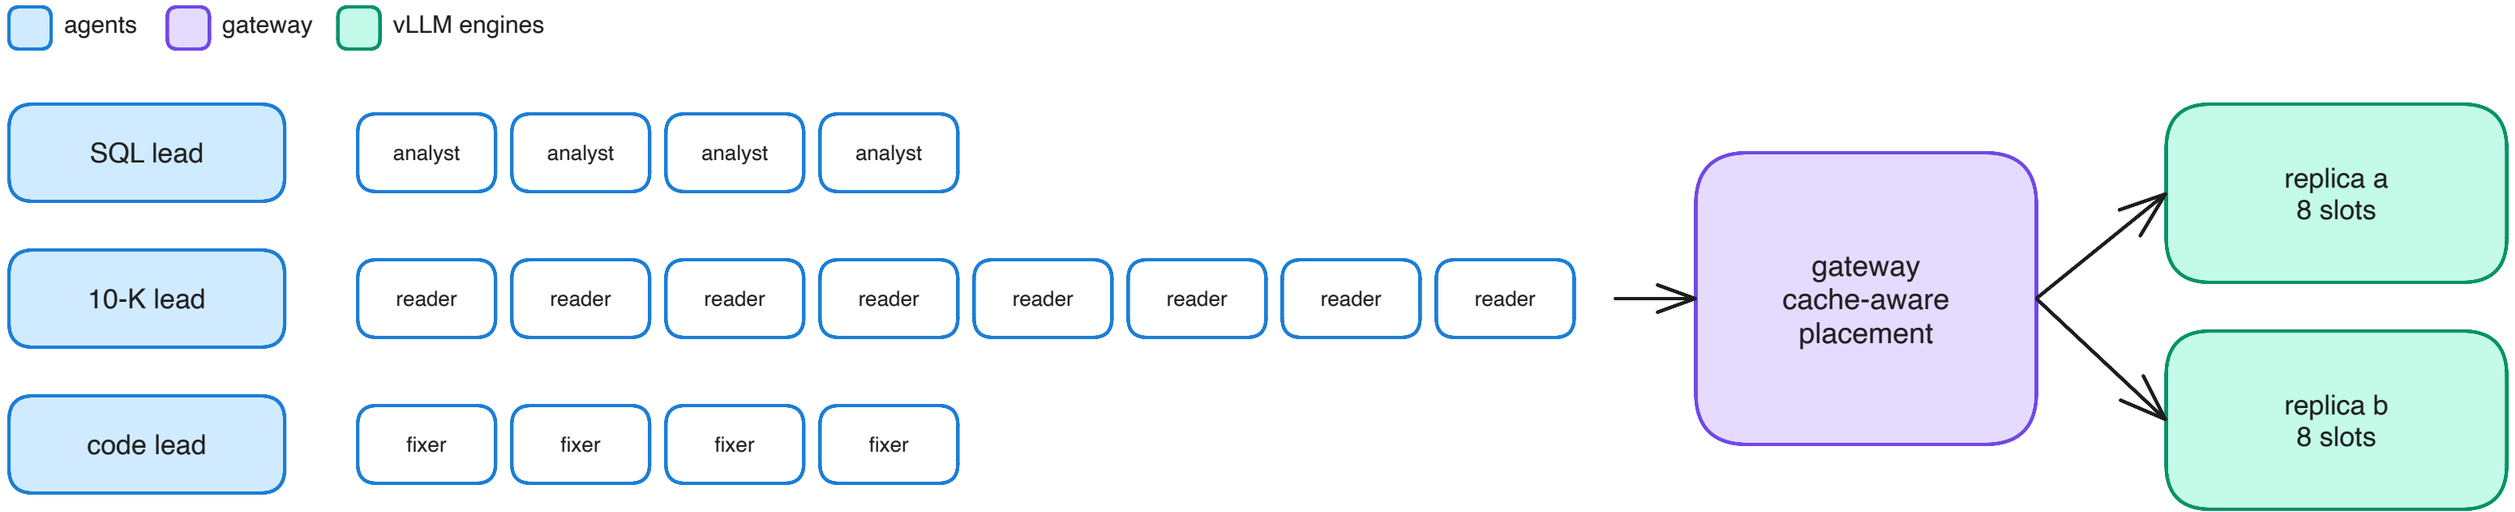

Leads orchestrate; at the peak, 16 subagent calls ran at once: every slot (8 readers + 4 + 4). Two readers share each ~10K-token filing, so the gateway keeps each pair on the replica that has it cached: 77% of prompt tokens came from cache.
Serving held up (607 calls, split 300 / 307, no preemption). The limits were the agent harness and the 9B's quality.