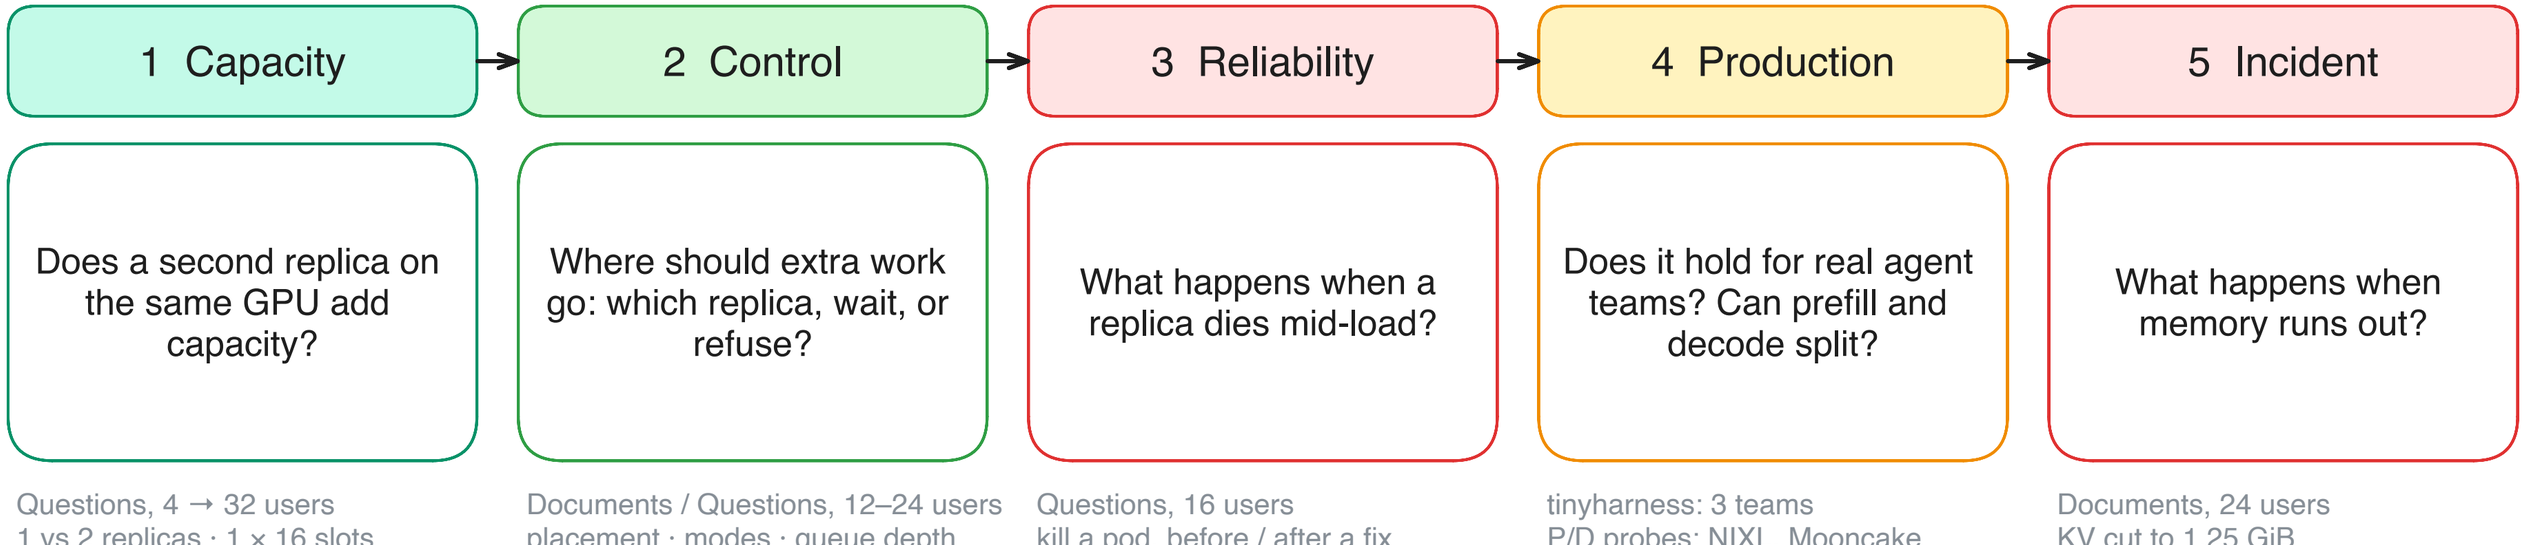
1 Capacity
2 Control
3 Reliability
4 Production
5 Incident
Does a second replica on the same GPU add capacity?
Where should extra work go: which replica, wait, or refuse?
What happens when a replica dies mid-load?
Does it hold for real agent teams? Can prefill and decode split?
What happens when memory runs out?
Questions, 4 → 32 users
1 vs 2 replicas · 1 × 16 slots
Documents / Questions, 12–24 users
placement · modes · queue depth
Questions, 16 users
kill a pod, before / after a fix
tinyharness: 3 teams
P/D probes: NIXL, Mooncake
Documents, 24 users
KV cut to 1.25 GiB